

# Intelligent Underwriting Workbench

An Agentic AI Co-Pilot for Commercial Risk Assessment

**LangChain**

**Ollama**

**Streamlit**

**Pydantic**

Presented by: Gurupreet Dhande, Khushi Dekate, Arnav Kalambe

# The Business Challenge

## Problem Statement

### Unstructured Data Overload

Commercial underwriting involves analyzing messy, unstructured documents:

- PDF Application Forms
- Loss Run Reports (Claims History)
- External Credit Reports
- Broker Emails & Notes

### Operational Impact

Manual processing leads to:

-  **Slow Turnaround:** Days to quote.
-  **Inconsistency:** Human error/bias.
-  **High Cost:** Experts doing data entry.

# The Solution: Agentic AI Workbench

Overview

A "**Human-in-the-Loop**" **Co-Pilot** that automates data extraction, analysis, and initial decisioning.



## Ingest

Reads messy PDFs and raw text instantly.



## Analyze

Parallel agents assess Claims, Credit, and Legal risks.



## Decide

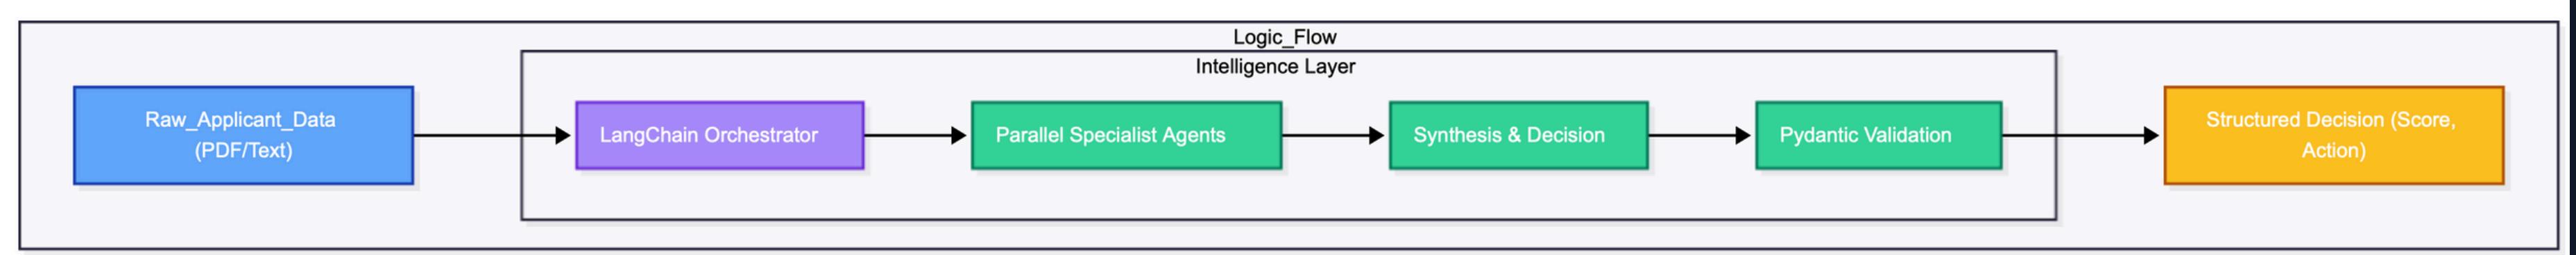
Generates a consistent Risk Score (0-10).



## Act

Drafts formal decision letters automatically.

# Detailed Components

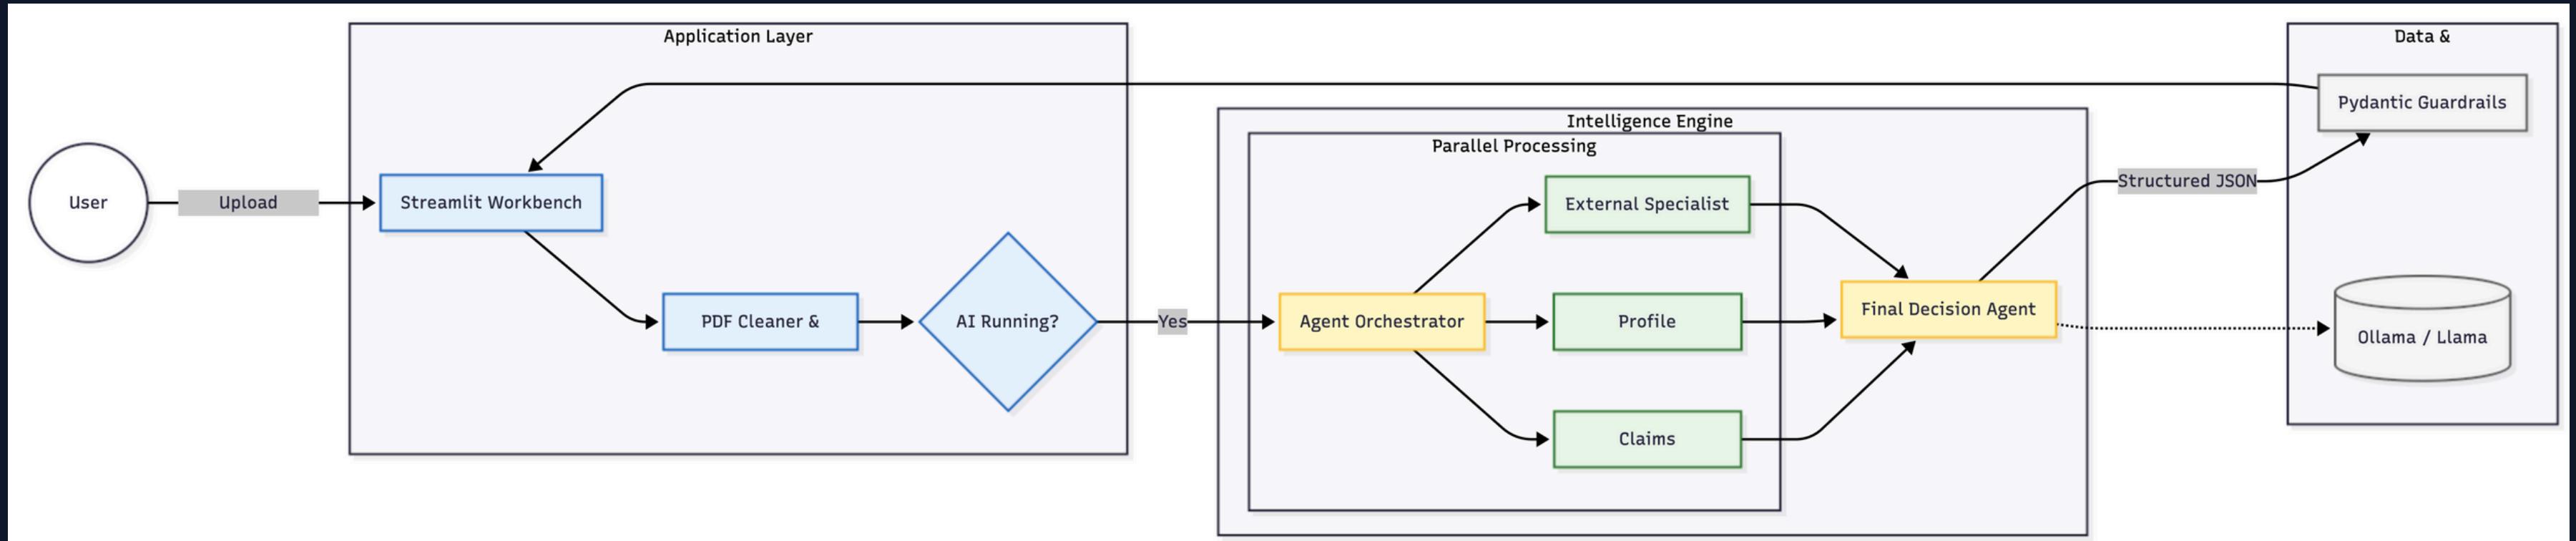


## Core Layers

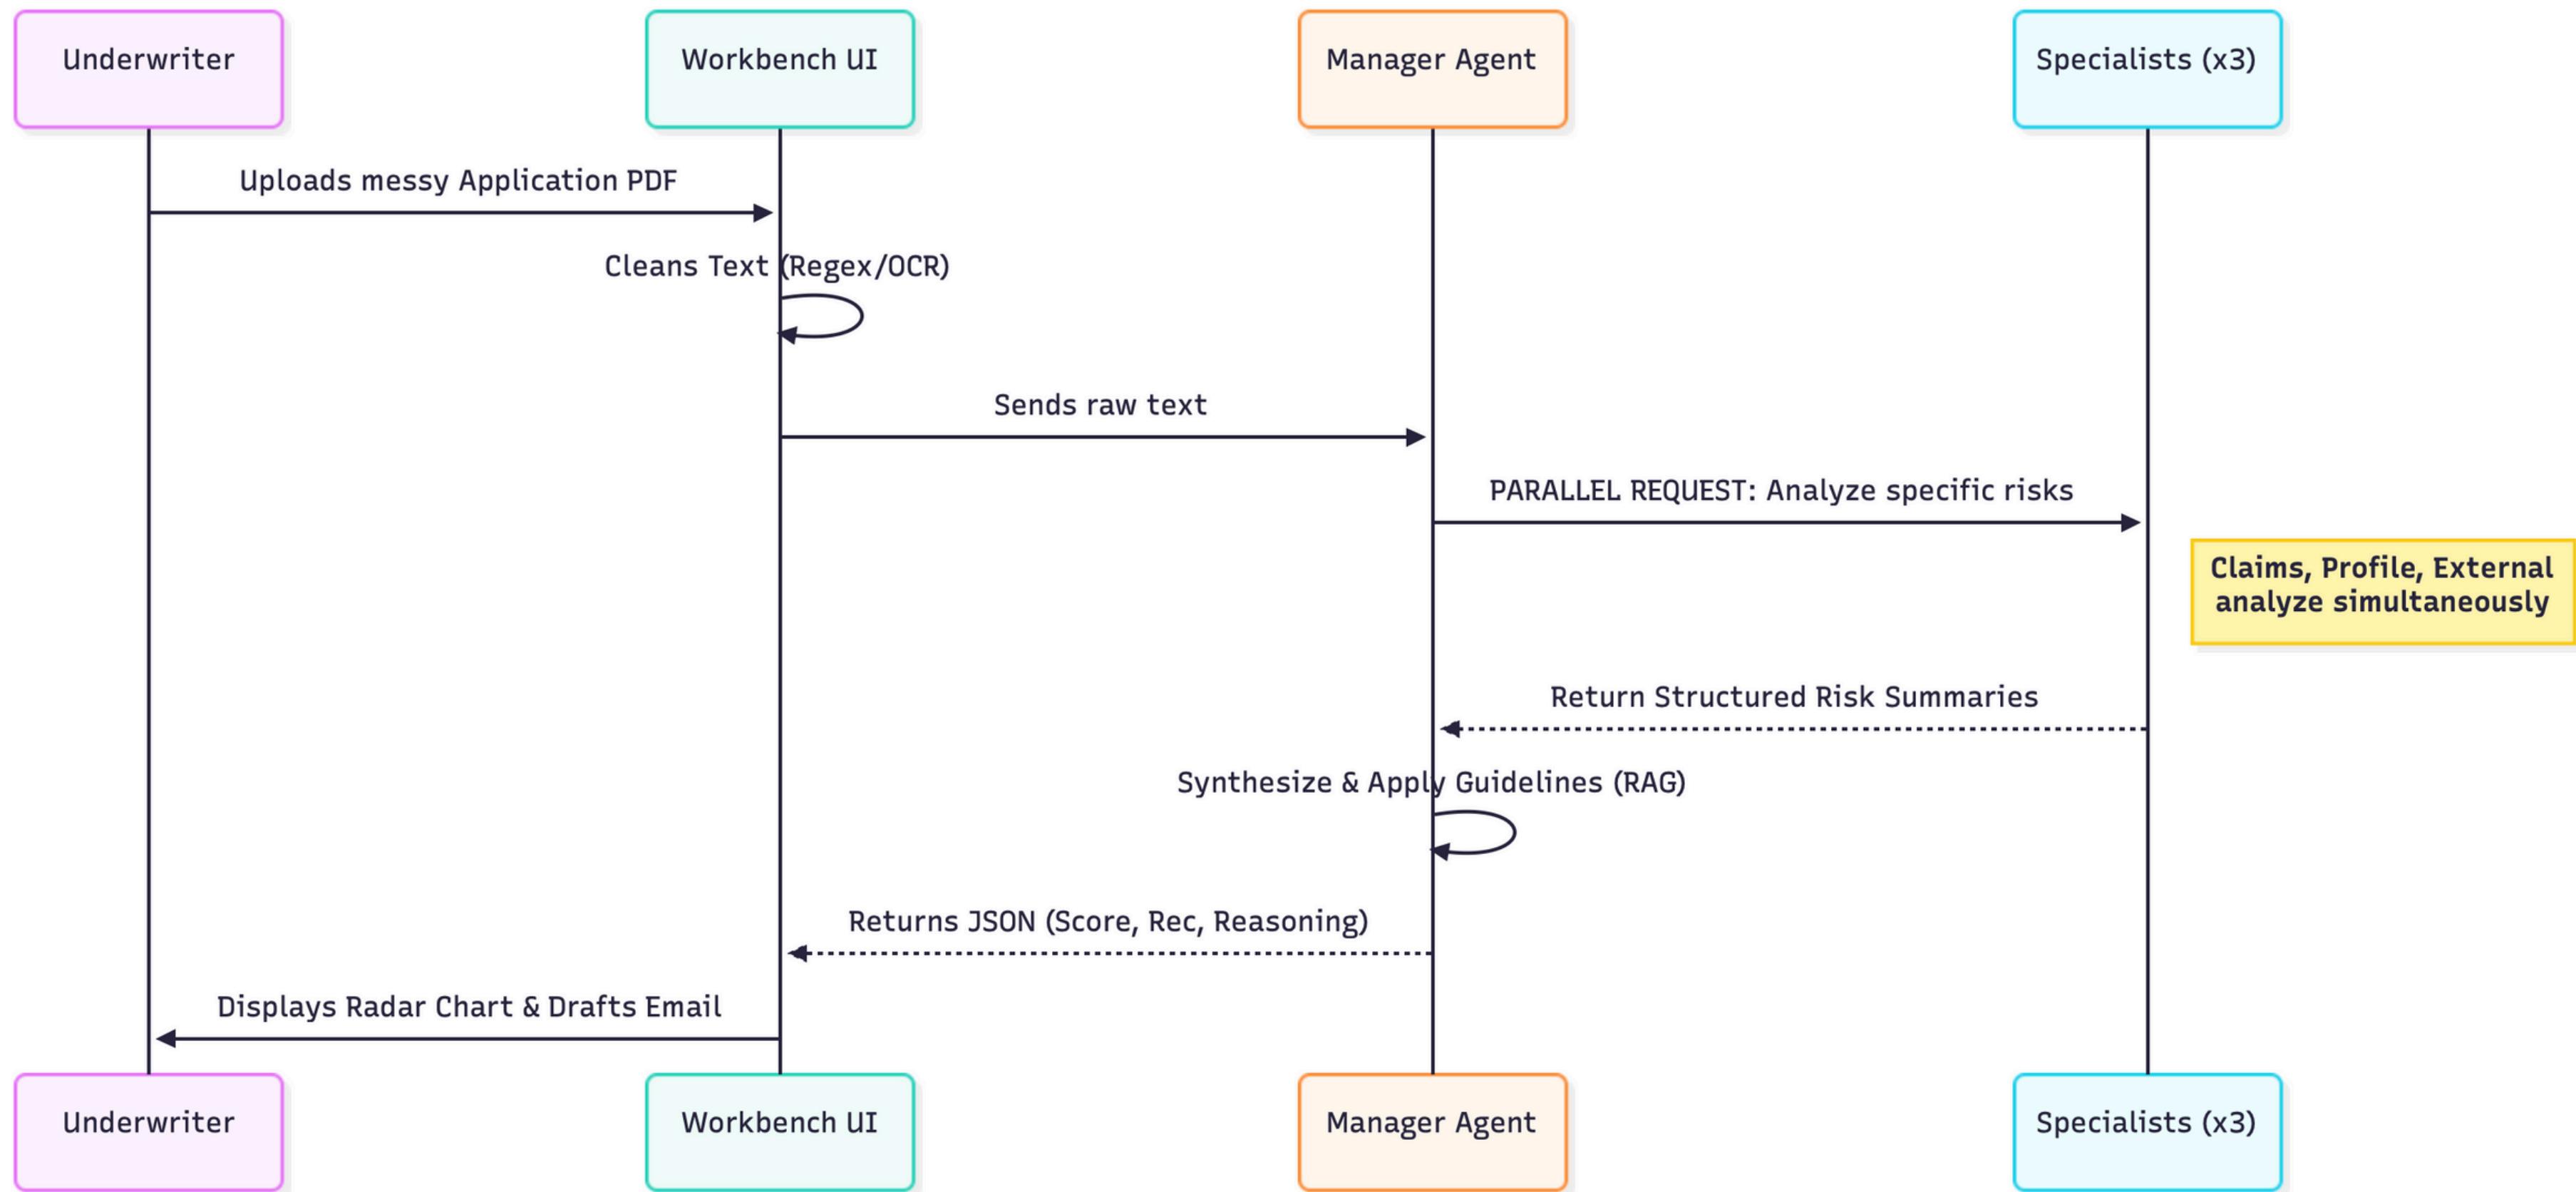
- **Ingress:** Nginx/Traefik Load Balancer.
- **Frontend:** Streamlit (Dashboard, Visuals).
- **Backend:** FastAPI (Async endpoints, Scaling).
- **Intelligence:** LangChain + Ollama + XGBoost.
- **Storage:** PostgreSQL (Audit), S3 (PDFs).

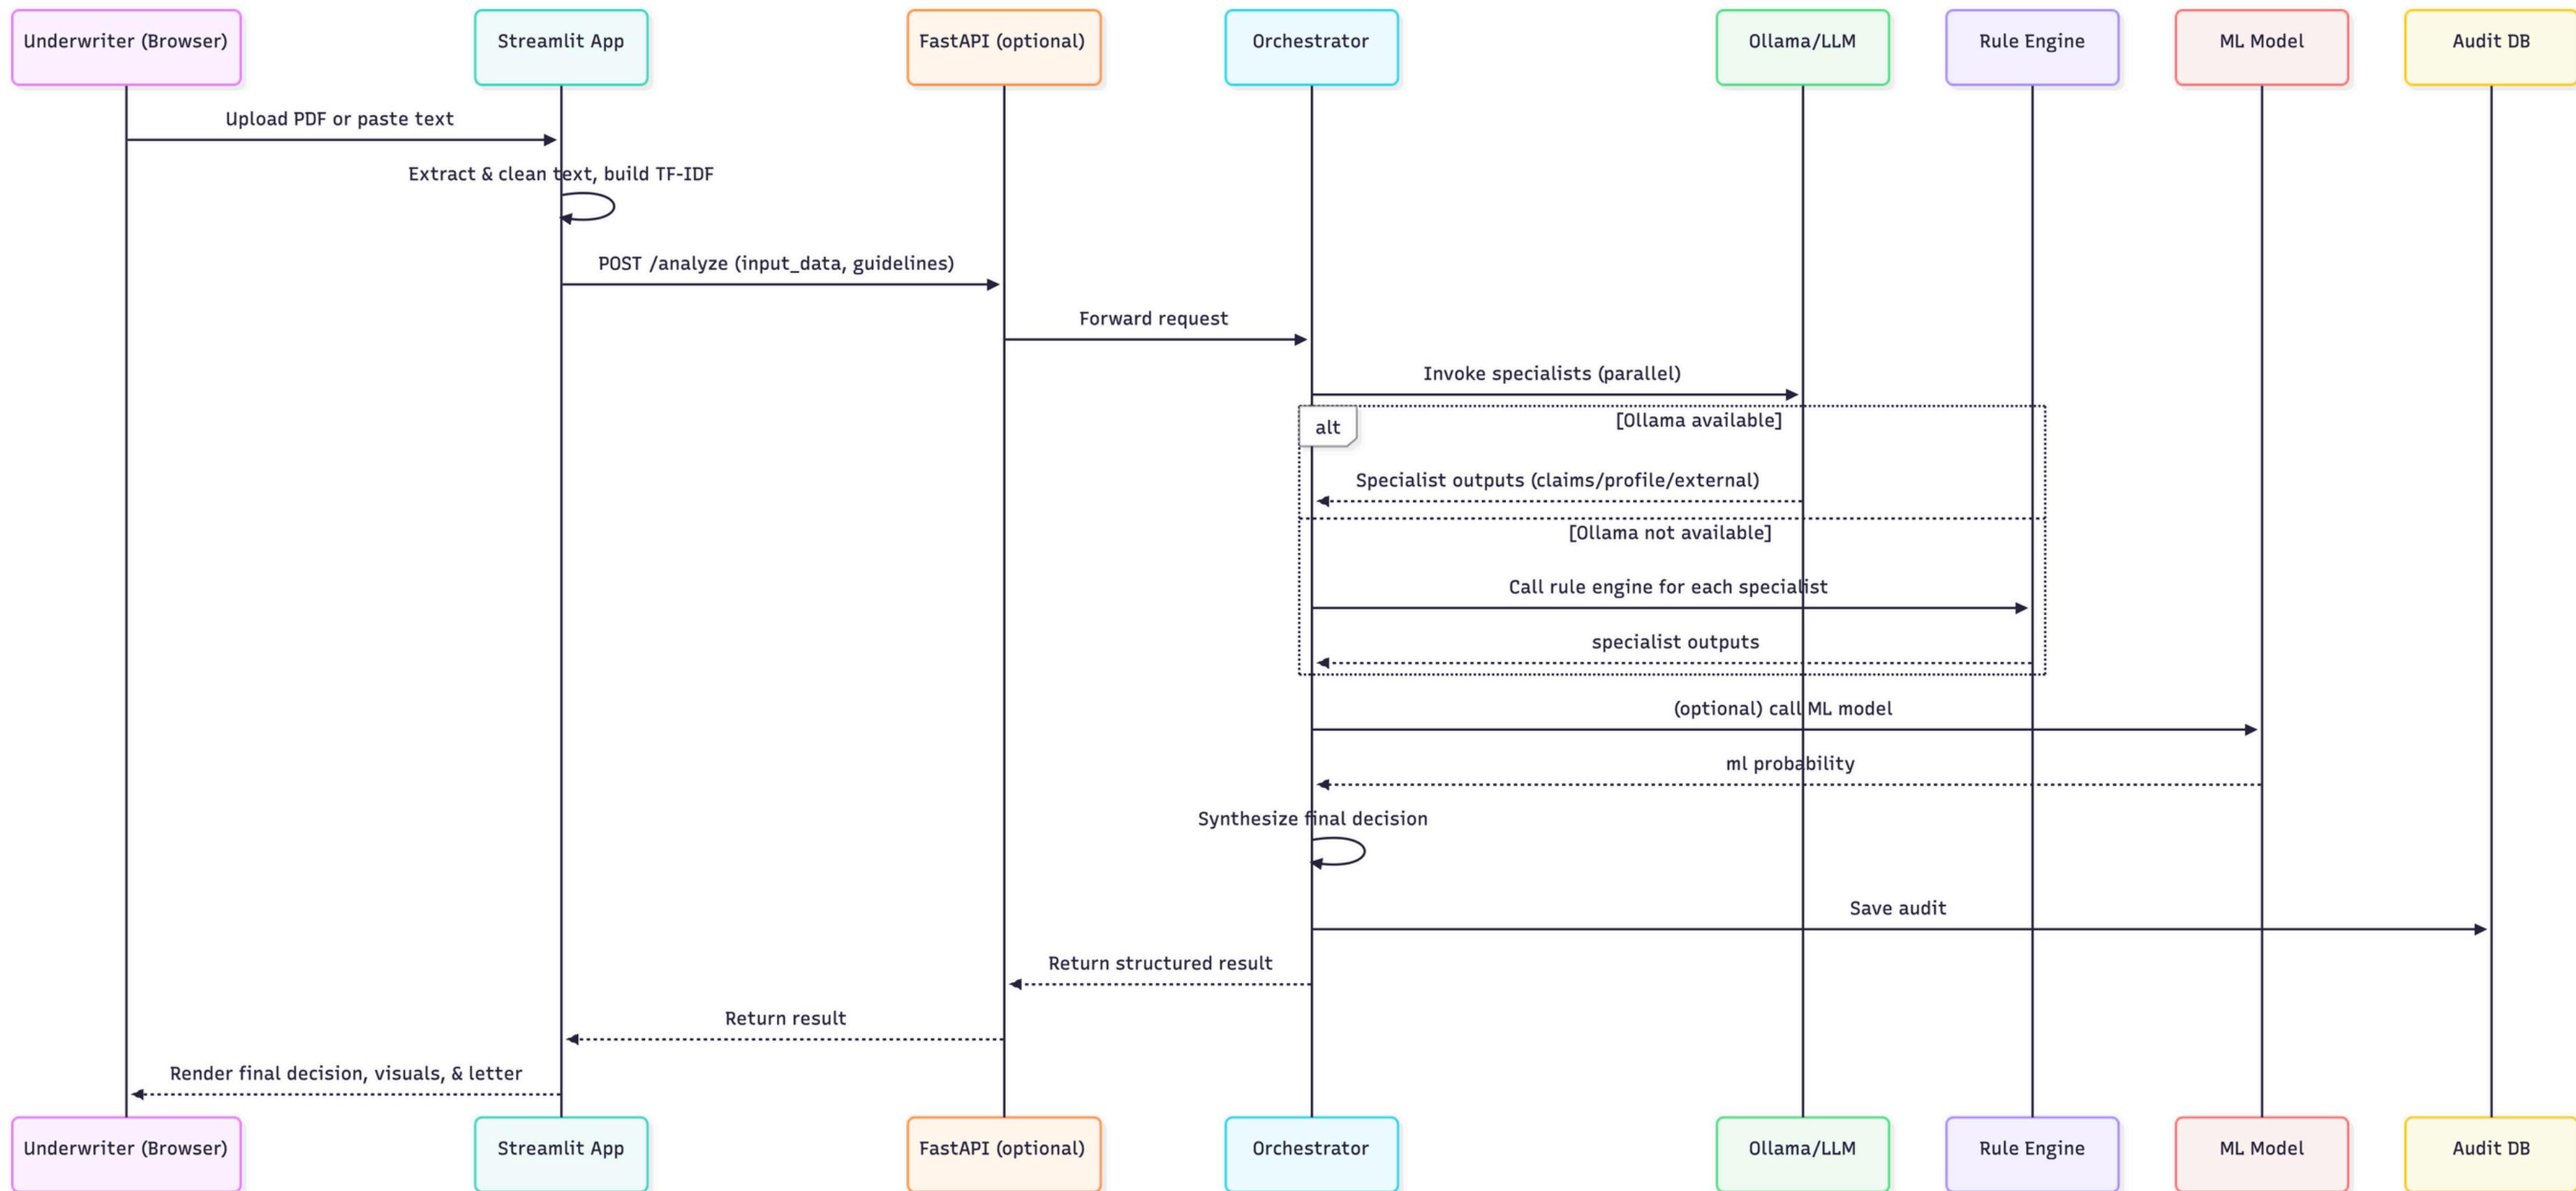
# System Architecture

High Level Design



The design of the Intelligent Underwriting Workbench, a robust, hybrid system. It focuses on demonstrating the Agentic Architecture, showcasing solutions to technical challenges like PDF parsing and JSON reliability, and concluding with the business impact of rapid, consistent risk assessment.





# The "Multi-Agent" Methodology

Parallel Processing

Why multiple agents? Because a single LLM gets overwhelmed. We use **Separation of Concerns**.

## 1. Claims Specialist

**Role:** Looks ONLY for loss history.

**Task:** Extract frequency, severity, and recency of claims.

```
{"score": 6, "summary":  
"Large fire claim in  
2021..."}
```

## 2. Profile Specialist

**Role:** Looks ONLY for business health.

**Task:** Analyze credit score, years in business, and industry risk.

```
{"score": 2, "summary":  
"Strong credit, 15yr  
tenure..."}
```

## 3. External Specialist

**Role:** Looks ONLY for legal/compliance.

**Task:** Find OSHA violations, lawsuits, or negative news.

```
{"score": 4, "summary":  
"Minor safety  
violations..."}
```

# Feature 1: Robust PDF Ingestion

Technical Deep Dive

## The Challenge

PDF extraction often results in "broken text" where lines split mid-sentence or tables lose structure.

*Example Raw Output:*

```
Appli-  
cant Name: Ac-  
me Const...
```

## The Solution (Regex Pipeline)

We built a custom cleaning pipeline in Python:

1. **Merge Hyphens:** Reconnects split words.
2. **Fix Newlines:** Aggressive regex `([^\n])\n([^\n])` to merge broken lines while keeping paragraphs.
3. **Standardize Bullets:** Converts confusing PDF symbols to standard text bullets.

# Feature 2: Structured Guardrails

Pydantic Integration

## Without Guardrails

LLMs are unpredictable. They might return:

- A long paragraph.
- Invalid JSON.
- A score of "High" instead of a number.

**Result:** The application crashes.

## With Pydantic & LangChain

We define a strict **Data Contract**:

```
class RiskOutput(BaseModel):  
    score: int = Field(ge=0, le=10)  
    decision: Literal["Approve", "Decline"]
```

**Result:** 100% Reliable, machine-readable output.

# Feature 3: Advanced Capabilities

RAG & Analytics

## Context-Aware RAG

The agent doesn't just guess; it reads **Underwriting Guidelines** (injected via sidebar) to make policy-compliant decisions.

## Risk Visualization

Used **Plotly** to generate interactive Radar Charts, visualizing the balance between Credit, Claims, and Legal risks.

## Explainability (XAI)

Includes a "Deep Dive" dashboard using **SHAP** concepts to show exactly *why* a score was given (e.g., "Credit score reduced risk by 10%").

# Deployment Strategy

---

## Production Scaling

- **Orchestration:** Kubernetes (K8s).
- **Async Tasks:** Redis + Celery for batch jobs.
- **Storage:** PostgreSQL (Audit) + S3 (Docs).

## Containerization

- **Frontend Container:** Streamlit
- **Backend Container:** FastAPI + Orchestrator
- **Model Container:** Ollama (GPU enabled)

# Demo Flow

# Why This Tech Stack?

Technology Choices

## Why Ollama / Llama 3?

**Privacy & Security.** Insurance data is sensitive. Cloud APIs (like GPT-4) send data off-premise. Ollama runs 100% locally, ensuring zero data leakage.

## Why LangChain?

**Orchestration.** Managing 4 distinct agents and passing data between them requires a robust framework. LangChain handles the "plumbing" effortlessly.

## Why Streamlit?

**Speed to Value.** Allowed us to build a full-stack, interactive dashboard in Python without needing a separate frontend team.

## Why Pydantic?

**Reliability.** In enterprise software, AI outputs must be predictable types, not random text.

# Impact & Results

Business Value

**90%**

Faster Triage

Reduced initial review time from  
30 mins to 30 seconds.

**100%**

Consistency

Eliminates human variability in  
risk assessment.

**0\$**

API Costs

Runs completely locally using  
Open Source models.

# Future Roadmap

Next Steps

## Phase 1 (Current)

- ✓ Multi-Agent Analysis
- ✓ PDF Ingestion
- ✓ Rule-Based Fallback

## Phase 2 (Next)

-  **Full RAG:** Connect to vector DB of 10,000+ past policies.
-  **Human Feedback Loop:** Allow underwriters to correct AI to improve model.

## Phase 3 (Scale)

-  **API Integration:** Connect directly to Policy Admin Systems.
-  **Multi-Modal:** Analyze photos of property damage.

**Thank You**